


The Cost Collapse

How quickly has the cost of accessing capable AI fallen?

API price per 1M tokens, major model releases 2020-2026



-99.25%

blended price of a capable token:
GPT-3 \$60.00 (2020) → GPT-5.6
Luna \$0.45 per 1M tokens (Jul 2026)

133×

cheaper in 6 years — before batch (-50%) and cached-input (-90%) discounts, which make it steeper

-80%


GPT-5.6 Luna price cut on Jul 30, 2026, three weeks after launch — verified via OpenAI + CNBC

Read: every dot is a model's list price at launch or reprice (log scale — each gridline is 10× cheaper). The floor falls relentlessly; premium flagships (GPT-5.5, Fable 5) ticked back up in 2026 even as any fixed capability level kept getting cheaper.

Intelligence per Dollar

How many times more capability does one dollar buy today?

Capability vs. blended cost — bubble size = context window



2,790×

more intelligence per dollar than GPT-3 — best-value model (Gemini 2.0 Flash, Feb 2025)

2,160×

GPT-5.6 Luna (Jul 2026): 93-point capability at \$0.45 per 1M tokens


≈300×

even the premium frontier tier (Grok 4.5, Claude Opus 5) buys ~300× more capability per dollar

Index method: capability = composite of GPQA-Diamond (normalized above the 25% chance floor) and SWE-bench Verified; pre-2023 models chained from MMLU at the GPT-4 overlap. Index = capability ÷ blended \$/1M tokens, GPT-3 (2020) = 1. Estimates flagged in the dataset; benchmark saturation makes recent gains conservative.

The Super Moore's Law

AI doesn't literally follow Moore's Law — but capability per dollar is currently compounding $\sim 3\times$ faster than Moore's pace.



1 –99%+

total cost decline: \$60 $\rightarrow$ \$0.45 per 1M blended tokens (133 $\times$); GPT-3-class capability $\sim 1,000\times$ cheaper (a16z)

2 $\sim 2,800\times$

improvement in capability per dollar since GPT-3 (best-value tier); $\approx 300\times$ at the premium frontier

3 $\sim 6-9$ mo


doubling time of AI capability per dollar — vs 24 months for the Moore's Law benchmark ($\approx 8\times$ over the same six years)

Corroboration: Epoch AI: constant-capability inference prices fell 9-900 $\times$ /yr (median $\sim 50\times$). a16z: $\sim 10\times$ /yr. Stanford HAI: GPT-3.5-class cost fell 280 $\times$ in 18 months. MIT FutureTech: 5-10 $\times$ /yr. Our 6-9-month doubling is the conservative end.

Appendix: The Open-Weight Catch-Up

Open-source models are the engine of the price collapse — and now trail the closed frontier by months, not years.

Open-weight models have nearly closed the capability gap — and set the price floor



~3-6 mo

capability lag of the best open-weight model behind the closed frontier in 2026 — down from ~a year in 2023

1/27th

DeepSeek R1's output price vs OpenAI o1 at launch (Jan 2025: \$2.19 vs \$60 per 1M) — the 'DeepSeek shock'

\$0.87

output price of open-weight DeepSeek V4-Pro (MIT license, 1M context) vs \$25-50 for closed flagships

Open-weight examples in the dataset: Meta Llama 2 / 3.1 405B / 4 · DeepSeek V2, V3, R1, V3.2, V4-Pro (MIT) · Mistral Mixtral & Large 2/3 (Apache 2.0) · Qwen 3 (Apache 2.0) · Kimi K2 / K3 (modified MIT — K3 is the largest open-weight model ever at 2.8T params). Open-weight releases triggered the May-2024 China price war and keep resetting the industry price floor.